

PropertyScope NSW

Release 0 / Technical report

An integrated NSW property-research workspace, built by five feature teams on a shared microservices foundation.

41026 Advanced Software Development

Assessment 1: Agentic AI Foundations, Microservices and DevOps

Submission	Group 20
Due	6 September 2026, 11:59 pm Sydney time
Repository	MattShelton04/41026ASDProject
Demonstration	Watch the published group demonstration
Showcase	Recorded and presented in Week 6 on 4 September 2026
Software baseline	7d5350d19023fb1e978e85127a72e3500a1556f3

Five integrated features. One shared application.

Property discovery / Market intelligence / Suburb analytics / Due diligence / Buyer journey

This report distinguishes implementation, retained execution evidence and remaining limitations. Repository evidence links are pinned to the software baseline above; report figures are regenerated from versioned Mermaid sources.

Contents

Select a chapter or page number to navigate. PDF bookmarks include every subsection.

Executive summary	3
1 Project overview	3
2 Agile project analysis and planning	4
3 Repository structure	14
4 Software architecture	15
5 AI mode and agentic workflow	22
6 Implementation summary	25
7 DevOps and GitHub Actions	27
8 Testing and validation evidence	29
9 Known issues and limitations	33
10 Contribution and commit record	33
11 Rubric traceability	35
Appendix A Evidence index	36
Appendix B Reproduction commands	36

Executive summary

PropertyScope NSW is an integrated research application for NSW property information. It brings property identity, recorded sales, suburb and crime context, planning and building evidence, and a buyer's own cases into one shared interface. Five independently owned feature slices run behind a containerised HTMX entry point and common visual system. Each feature has a frontend, backend API, owned database service and visible CRUD workflow. Backends call the shared AI mode over HTTP and do not import another student's implementation or open another feature's database.

The Release 0 baseline contains all five enabled features and 21 Compose services. The repository's canonical quality gate passed locally on 4 September 2026. Each student workflow also has a retained successful GitHub Actions run, and those runs build containers and exercise an isolated feature stack. Features 2, 3 and 5 use SQLite behind an internal database API. Features 1 and 4 use PostgreSQL/PostGIS. Each feature retains database ownership behind its private database API.

AI mode implements a durable Plan, Act, Observe and Adapt state machine with versioned planner and adapter prompts, schema-validated feature tools, bounded retries and human review for protected actions. OpenAI is the approved Release 0 provider, recorded in the registered feature scope. The OpenAI Responses API routes planning to GPT-5.6 Luna and adaptation to GPT-5.6 Terra. The team selected this API for low-cost experimentation and its account's data-sharing token allowance (Section 5.1). Deterministic CRUD and evidence views continue to work without a model credential.

The recorded demonstration and Week 6 presentation are complete. In accordance with the 30 August showcase clarification, the video covers the integrated application, per-feature AI paths, deployment and CI/CD; the agent-loop execution evidence is retained in this report instead of being repeated in the video.

1 Project overview

1.1 Problem and product purpose

NSW property research is fragmented across government datasets, property identifiers, recorded sale history, suburb indicators, planning constraints and personal notes. PropertyScope gives a user one place to organise that evidence while keeping source, coverage, freshness and uncertainty visible. It supports research and comparison. It does not provide valuation, legal, safety, lending, conveyancing or buy and sell advice.

Release 0 establishes the software foundation: integrated microservices, owned persistence, visible CRUD, AI mode, a shared agent loop, Docker Compose and continuous integration. Later-release MCP, RAG, multi-agent and cloud deployment work is outside this report.

1.2 Team and feature allocation

Student	Name	Student ID	Feature	Release 0 result
1	Matthew Shelton	24763373	Data Platform and Property Discovery	Enabled and integrated
2	Burhan Naeem	24764134	Property Sales Explorer and Market Cases	Enabled and integrated
3	James Huang	24970865	Suburb Crime and Liveability Analytics	Enabled and integrated
4	Michael White	24846267	Site Planning and Building Due Diligence	Enabled and integrated
5	Derek Song	24833978	Buyer Journey and Agent Workspace	Enabled and integrated

The approved ownership boundary is recorded in [registered feature scope](#). Allocation does not permit one student service to import another student's production code. Cross-feature integration occurs through public HTTP APIs and versioned data products.

1.3 Release 0 acceptance map

Assessed capability	Implemented evidence
Shared setup	Locked Python 3.12 workspace, five student folders, shared packages, AI services, Docker Compose, common CSS and HTMX entry point
Service implementation	Five enabled frontend, backend and database-service slices; 21 services in the generated release profile
AI mode	Shared durable run API, provider registry, versioned prompts and 25 registered feature tools with protected writes review gated
Agentic workflow	Plan, Act, Observe and Adapt state machine, events, evidence, recovery and review states
DevOps	Integration CI plus student -1.yml through student -5.yml, all with retained successful runs
Working software	Visible CRUD in every feature and deterministic offline paths
Evidence	Local quality results, CI URLs, container smoke results, screenshots, run IDs, commit and contribution records

2 Agile project analysis and planning

2.1 Working method

The team used short-lived feature branches and pull requests into main. The repository moved from the shared platform and Feature 1 baseline to five integrated slices between 26 July and 4 September. Every merge triggered Integration CI, while path-filtered student workflows exercised the affected slice. Shared boundaries were expressed as contracts, feature manifests and executable architecture validators so integration decisions were checked by code rather than meeting notes alone.

The working cycle was:

1. define or refine a feature-level acceptance outcome;
2. implement the smallest vertical change across owned frontend, API, persistence and tests;

3. open a pull request and use CI evidence to find integration defects;
4. correct the defect without weakening the gate; and
5. merge only when the relevant checks passed.

This pattern is visible in the delivery history. Feature 3, for example, followed its initial vertical slice with focused pull requests for ingestion, filters, comparisons and cross-platform CI. Feature 4 similarly separated initial integration, CRUD, evidence, mapping, AI and workflow work.

2.2 Release backlog and completion

ID	Backlog outcome	Owner	Acceptance evidence	Result
R0 01	Shared repository contracts and product shell	Shared and Student 1	Locked workspace, shared route, CSS and architecture checks	Complete
R0 02	Durable AI mode and agent loop	Shared and Student 1	Run APIs, v7 prompts, tool catalogues and state-machine tests	Complete
R0 03	Property discovery and governed data platform	Student 1	CRUD, acquisition, publication, property search and CI	Complete
R0 04	Sales explorer and market cases	Student 2	Case CRUD, deterministic summaries, AI tools and CI	Complete
R0 05	Suburb crime and liveability analytics	Student 3	Search, comparison CRUD, import boundary, AI tools and CI	Complete
R0 06	Site planning and building due diligence	Student 4	Review CRUD, evidence states, AI questions and CI	Complete
R0 07	Buyer journey workspace	Student 5	Case, shortlist, note and task CRUD, evidence and AI summary	Complete
R0 08	One generated Compose application	Team	Five enabled manifests and 21-service release profile	Complete
R0 09	Five student workflows	Each owner	Successful run URL and SHA for every workflow	Complete
R0 10	Report video and Week 6 presentation	Team	This report and completed group demonstration	Complete subject to accessible video link

2.3 Common requirements

Every slice contributes the following functional requirements: an independently deployed frontend; a versioned backend API; an owned database API and schema; visible Create, Read, Update and Delete; shared-shell navigation and visual tokens; a bounded AI-mode path; health endpoints; and deterministic tests. The database API is the only process permitted to open its database file or use database credentials.

Quality area	Requirement	Evidence
Reproducibility	Python 3.12 and all workspace dependencies reproduce from <code>uv.lock</code>	Local and Integration CI locked sync
Isolation	No cross-student imports, shared database access or foreign volume mounts	<code>scripts/validate_architecture.py</code>
Resilience	CRUD and evidence paths remain available without a model provider	Offline tests and feature smoke jobs

Quality area	Requirement	Evidence
Validation	Untrusted identifiers, payloads, pagination and tool arguments are bounded	API and domain tests; JSON schema tool catalogues
Observability	Requests and agent runs retain safe status, events and evidence	Run APIs, request IDs and activity views
Accessibility	Keyboard controls, labels, visible focus and non-colour state cues	Frontend behaviour tests and shared UI audit
Maintainability	Typed boundaries, application factories, migrations and deterministic tests	Ruff, Mypy, pytest and Node checks
Data integrity	One owning database service and idempotent migrations	Compose volumes, internal APIs and migration tests

2.4 Individual feature plans and risks

Student 1 Data Platform and Property Discovery

The feature registers official sources, runs bounded acquisition, preserves immutable artefacts, validates candidates, requires human publication review and exposes accepted property evidence. Its frontend also provides Source CRUD and property search. The runner and backend use the private database API; only the database API and loader hold PostgreSQL credentials.

Planning area	Release 0 design
Functional backlog	Source and job CRUD; acquisition monitoring and cancellation; release review; property search; evidence inspection; AI failure diagnosis
Non functional focus	Provenance, bounded source-scale processing, atomic publication, cancellation, idempotency and deterministic offline operation
Conceptual model	Sources produce runs and artefacts; runs produce candidate releases; accepted releases populate a property registry and attributed warehouses
Logical model	ops source, job, run, task, artefact, release, quality and publication records; registry property identity; warehouse G-NAF, PSI, BOCSAR, school and SEIFA facts; serving accepted generations and coverage
Physical model	PostgreSQL 16 with PostGIS, schema migrations, private database API, separate loader and named database and artefact volumes
Principal risks	Publisher scale or change, incomplete evidence presented as certain and model failure
Controls	Content hashes, typed canonical artefacts, explicit coverage, safe capacity limits, human review and accepted-generation pointers

The ERDs show selected relationships in the **current Release 0 schema after all migrations through 049**, rather than the initial schema alone. The [complete migration set](#) and [migration runner](#) define the physical schema; SQL files apply in filename order. Figure 1 covers operational lineage, including the release terminal-reason field introduced in migration 031.




Figure 1 Feature 1 operational lineage after migrations through 049

Feature 1 / Publication and data-product schema

The current schema separates immutable source evidence, reviewed release activation, downstream delivery and accepted data products. The following ERD complements the operational lineage above. Table names are shortened for readability: activation and delivery tables belong to ops, accepted_generation belongs to serving, and seifa_sal belongs to warehouse. Only selected keys and relationships are shown; the complete migration set remains the physical-schema reference.

Schema change	Current effect	Migration evidence
Run-scoped artefact lineage	Repeated content retains a separate lineage record for each ingestion run	004
Asynchronous release activation	Publication receipts and leased loader work are distinct durable records	026 activation
Terminal releases and loader progress	Abandoned releases retain a terminal reason; task and loader progress is persisted	031
Durable consumer delivery	Consumer operations, activation attempts and delivery aliases preserve retry identity	040 , 041 , 042
ABS SEIFA 2021	Locality measures retain release, artefact and ingestion provenance	048
G-NAF identity lookup	Additional indexes support exact address matching and historical identity anchors	049

Migration 049 changes indexes rather than adding an entity. Its effect belongs in the physical-design record above; the ERDs retain the resulting table relationships.

Feature 1 / Current publication and accepted-data relationships




Figure 2 Feature 1 publication delivery and SEIFA relationships after migrations through 049

Student 2 Property Sales Explorer and Market Cases

The feature stores market cases against externally owned property references, imports compatible sale observations, calculates deterministic summaries and asks AI mode to explain the recorded facts. Case CRUD uses optimistic versions and returns HTTP 409 for stale changes.

Planning area	Release 0 design
Functional backlog	Select a known property, create and edit a case, filter attributed sales, show exclusions and source releases, delete a case and request an AI explanation
Non functional focus	Deterministic calculations, bounded imports, provenance, version conflicts and offline use
Conceptual model	A market case selects a property and filters reusable sale observations; observations are not owned by a case; a summary is derived rather than stored as model output
Logical model	market_case and sale_observation; unique source business key, revision and release; property and contract-date index
Physical model	SQLite volume owned by f2-db-api; Nginx frontend and Flask backend and database APIs
Principal risks	Synthetic demonstration data, weak or missing price records, unavailable Feature 1 validation and model overreach
Controls	Explicit exclusions, integer median, source metadata, property_validation_state, read-only tools and no valuation or recommendation language

Selected keys and relationships; the physical schema defines the remaining fields, checks and indexes. Dashed associations labelled "no FK" are query relationships, not database constraints. [Physical schema](#).




Figure 3 Feature 2 entity relationships and selected physical keys

Student 3 Suburb Crime and Liveability Analytics

The feature provides suburb search, filters, maps, amenities, factual comparison views, saved comparison CRUD and optional AI explanation. It also implements a durable consumer for accepted Feature 1 crime, school and SEIFA products, although those official products were not activated in the audited Release 0 runtime.

Planning area	Release 0 design
Functional backlog	Search and filter suburbs; show context and amenities; compare two to five suburbs; save, edit and delete comparisons; request bounded AI analysis
Non functional focus	Zero versus missing semantics, accessible chart tables, import integrity, bounded distances, recovery and deterministic degradation
Conceptual model	A suburb has an overview, many indicators and many amenities; a saved comparison references multiple locality and indicator keys
Logical model	suburb_info, suburb_overview, suburb_indicators, suburb_amenity, user_suburbs; separate durable import, delivery and current-release records
Physical model	SQLite volume opened only by f3-database; dependency-free WSGI backend and Nginx frontend
Principal risks	Partial fixture mistaken for current official data, incompatible periods or units, boundary ambiguity and unavailable upstream release artefacts
Controls	Source labels, count and rate separation, null preservation, exact manifest and hash checks, atomic current pointers and no safety or desirability ranking

Selected keys and relationships; the physical schema defines the remaining fields, checks and indexes. Dashed associations labelled "no FK" are query relationships, not database constraints. [Physical schema](#).




Figure 4 Feature 3 entity relationships and selected physical keys

Student 4 Site Planning and Building Due Diligence

The feature creates and manages site reviews, combines planning, environmental, strata and building observations, shows explicit evidence states and generates questions for a qualified professional. The map is a deterministic visualisation around Feature 1 coordinates; it is not a parcel-accurate spatial intersection.

Planning area	Release 0 design
Functional backlog	Create, read, edit and delete site reviews; change checklist, status and notes; inspect attributed evidence; view a map; generate professional-verification questions
Non functional focus	Visible uncertainty, bounded GeoJSON, database isolation, provider degradation and safe non-advisory language
Conceptual model	Site reviews and observation collections share a property reference; observations are not children of a review
Logical model	site_review, constraint_observation, building_observation with evidence-state, source and confidence fields
Physical model	PostgreSQL/PostGIS behind f4-db-api, Nginx frontend and Flask backend
Principal risks	Synthetic evidence mistaken for official parcel evidence and model treated as compliance advice
Controls	Evidence badges and sources, explicit professional verification and read-only AI tools

Selected keys and relationships; the physical schema defines the remaining fields, checks and indexes. Dashed associations labelled "no FK" are query relationships, not database constraints. [Physical schema](#).




Figure 5 Feature 4 entity relationships and selected physical keys

Student 5 Buyer Journey and Agent Workspace

The feature manages buyer cases, shortlisted properties, notes and tasks, composes bounded evidence from other public feature APIs and requests a short evidence-aware AI case summary. Optimistic concurrency protects every mutable record.

Planning area	Release 0 design
Functional backlog	Buyer-case CRUD; shortlist CRUD; stage, rating and priority; note CRUD; task CRUD and completion; evidence refresh; AI summary
Non functional focus	Owner scope, version conflicts, safe external evidence, bounded context, request correlation, accessibility and persistence across restarts
Conceptual model	One buyer case owns many shortlist properties, notes and tasks; a note or task may refer to one property in the same case
Logical model	buyer_case, case_property, case_note, case_task; cascading case deletion, same-case triggers and unique shortlist entries
Physical model	SQLite volume owned by f5-db-api, Nginx frontend and Flask backend and database APIs
Principal risks	Incomplete cross-feature evidence, stale writes and model-generated unsupported action
Controls	Explicit availability states, demo-owner scope, positive versions and HTTP 409, source references, short read-only tool context and human decision warning

Selected keys and relationships; the physical schema defines the remaining fields, checks and indexes. [Physical schema](#).




Figure 6 Feature 5 entity relationships and selected physical keys

3 Repository structure

The monorepo keeps shared contracts domain-neutral and gives each student a vertical slice. Generated deployment files expose only validated enabled features.

<code>.github/workflows/</code>	Integration CI and student-1.yml through student-5.yml
<code>ai-services/</code>	Agent core and shared AI-mode service
<code>deployment/</code>	Enabled-feature registry and generated Compose overlay
<code>docs/</code>	Architecture decisions release evidence and report sources
<code>shared/</code>	Contracts consumer protocol testkit and shared frontend
<code>student-1/ to student-5/</code>	Independently owned frontend backend database tests and manifest
<code>scripts/</code>	Quality deployment fixture and operator commands
<code>docker-compose.yml</code>	Shared base service model
<code>docker-compose.dev.yml</code>	Reloadable local development overlay

The boundary validator enforces the key ownership rules: student services may import shared contracts and tests may import shared testkit, but a student service may not import another student or agent_core. AI mode may call allowlisted feature endpoints, never a student database. Each database volume is mounted only by its owner.

4 Software architecture

4.1 Individual software architecture

Feature 1 / Data Platform and Property Discovery

Feature 1 is deliberately more than a frontend, backend and database. The credential-free runner claims acquisition work through the backend's worker HTTP contract and is the only service that writes content-addressed source artefacts. The DB API owns request-time persistence access, while the separate credential-owning loader performs durable imports and accepted-generation activation. All three services read the artefact volume, but only the runner writes it; only PostgreSQL mounts the database volume. Publication remains a reviewed backend action and accepted releases are delivered to downstream consumers through idempotent HTTP contracts.




Figure 7 Feature 1 runtime trust boundaries acquisition loading AI callbacks and publication

The Mermaid source for every architecture, state and pipeline figure is retained under docs/reports/diagrams/release-0. The rendered PNG files are the PDF-compatible derivatives, not independent drawings.

Feature 2 / Market cases and recorded sales

Market-case CRUD crosses the frontend, backend and private database API. The backend validates property references through Feature 1 and derives sales statistics deterministically. Two read-only tools expose case context and summary facts to AI mode. The SQLite file remains exclusively with f2-db-api.




Figure 8 Feature 2 frontend backend persistence and AI boundaries

Feature 3 / Suburb and liveability analytics

The backend owns query, comparison, nearby-place and assistant routes plus the fenced import worker. Only f3-database opens SQLite. Feature 1 products arrive over validated HTTP delivery contracts; the retained Release 0 demonstration uses labelled fixture data. Three read-only tools bound the assistant context.




Figure 9 Feature 3 frontend backend persistence and AI boundaries

Feature 4 / Site and building due diligence

The backend validates properties through Feature 1, manages reviews and prepares bounded evidence and map responses. The private database API alone receives PostgreSQL credentials; f4-postgres owns the database volume. Two read-only tools support professional-verification questions.




Figure 10 Feature 4 frontend backend persistence and AI boundaries

Feature 5 / Buyer journey workspace

The backend manages cases, shortlist entries, notes and tasks through its private SQLite API. Public Feature 1, 2 and 4 APIs supply bounded research evidence. Feature 3 is reported as unavailable in this baseline. Four owner-scoped, read-only tools support the AI summary; version checks protect mutable records.




Figure 11 Feature 5 frontend backend persistence and AI boundaries

4.2 Integrated architecture

The browser normally enters through shared - frontend on port 5100. The shared edge loads a same-origin HTMX feature directory and proxies feature routes so users remain within one product shell; per-feature frontend ports remain available for local development. Backends call only their owned database service and the shared AI mode over private Compose networking. AI mode calls back to allowlisted feature-owned HTTP tools and never opens a feature database. Feature 1 publishes accepted dataset releases; downstream features consume those through supported HTTP contracts.




Figure 12 Integrated Release 0 software architecture

The integrated architecture contains additional workers only where durable background processing is required. Feature 1 separates acquisition from database loading so the runner never receives database credentials. Feature 3 performs fenced consumer imports through its backend worker. Feature 5 composes data at request time from public APIs and keeps the result's availability and limitations.

4.3 Docker Compose architecture

deployment/features.yaml enables all five manifests. The generator creates deployment/enabled-features.compose.yaml and the shared route projection. The resulting Release 0 profile contains 21 services and seven named volumes.




Figure 13 Docker Compose service groups

Group	Services	Owned durable volume
Shared	shared-frontend, shared-ai-mode	shared-ai-mode-state
Student 1	f1-frontend, f1-backend, f1-runner, f1-db-api, f1-db-loader, f1-postgres	f1-postgres-data, f1-artifacts
Student 2	f2-frontend, f2-backend, f2-db-api	f2-sqlite-data
Student 3	f3-frontend, f3-backend, f3-database	f3-suburb-data
Student 4	f4-frontend, f4-backend, f4-db-api, f4-postgres	f4-postgres-data
Student 5	f5-frontend, f5-backend, f5-db-api	f5-sqlite-data

The canonical local deployment command is:

```
uv run scripts/dev.py stack up
```

Compose configuration was validated on 4 September with all five features enabled. Docker Desktop was unavailable during the final report edit, so this report relies on the retained successful student workflow container jobs for runtime startup, health, seed and smoke evidence rather than claiming a new local Docker run.

5 AI mode and agentic workflow

5.1 AI mode design

Shared AI mode owns provider selection and the agent lifecycle. A feature backend submits an objective, trusted context, prompt-set version, limits and a feature-owned tool allowlist. AI mode validates the request, persists it to its SQLite run store and executes the bounded state machine. Feature tools are ordinary HTTP endpoints with validated JSON schemas. Protected writes do not run automatically; they require a separate reviewed action.

Provider decision: OpenAI preferred; Gemini supported. The approved [registered feature scope](#) records OpenAI with GPT-5.6 Luna and GPT-5.6 Terra. The remote-standard.v1 profile uses Luna for planning and Terra for adaptation through the Responses API. This is the team's settled Release 0 selection under the approved registration, and replaces the brief's example Ollama runtime. No feature selects a concrete model itself.

The team confirms that its data-sharing opt-in provides **2.5 million complimentary tokens per day for its Luna/Terra use**. That account allowance makes repeated prompt refinement, bounded agent runs and a five-student demonstration practical at low cost, without local model weights or GPU setup. The allowance is an account-specific planning assumption supplied by the team, not a universal or per-model entitlement. OpenAI's [data controls documentation](#) explains that API data is not used for model improvement unless the customer explicitly opts in.

The demonstration uses public-source evidence and labelled fixtures; credentials remain in a file secret outside the report and repository. Calls retain bounded iteration, tool and time budgets. Gemini API is also supported through the explicit gemini-development.v1 profile. Its [free tier for eligible models](#) provides another low-cost option for development and experimentation, subject to the selected model's rate limits. The Gemini adapter uses Google's OpenAI-compatible Chat Completions and Models endpoints; it is configured as a separate provider mode because these endpoints do not implement the Responses API. The [provider operations guide](#) documents this configuration. OpenAI remains the preferred default, while the Gemini run in Section 5.4 provides retained evidence that the alternative provider works through the same feature and agent boundaries.

5.2 Plan Act Observe Adapt




Figure 14 Bounded Plan Act Observe Adapt workflow

Phase	Durable behaviour	Evidence retained
Plan	Model selects only allowlisted tools and valid arguments within iteration and time limits	Prompt ID and version, provider, model, plan status
Act	Agent core validates the call and invokes the feature HTTP endpoint	Tool name, call ID, safe arguments, timing and HTTP outcome
Observe	Tool result is validated and stored as evidence	Result summary, evidence reference and failure classification
Adapt	Model integrates observations and either finishes, replans or requests review	Final findings, limitations, recommendation and next state

The diagram summarises the main execution path. The [normative state machine](#) specifies every failure, cancellation and recovery transition. The run can terminate as succeeded, failed or cancelled, or pause at review required. Idempotency, bounded recovery and durable checkpoints prevent an interrupted process from silently repeating a protected action.

5.3 Prompt engineering and context management

The shared prompt set is stored under `ai-services/ai-mode/src/ai_mode/prompt_assets`. Planner and adapter templates have matching metadata and immutable versions from v1 to v7. Version 7 added typed provenance and deterministic parallel read-only stages to the planner, then taught the adapter to report ordered results from sequential and parallel execution. Runs retain the exact prompt set and model profile used.

Contributor	Prompt or context asset	Bounded context and safety rule
Student 1	Shared planner and adapter v1 to v7 plus data-platform tool catalogue	Typed provenance, recorded coverage, no automatic publication and explicit unknown states
Student 2	Market-case objective in backend API and two-tool catalogue	Trusted case ID only; deterministic facts; no valuation, forecast or buy recommendation
Student 3	General and saved-comparison objectives plus three-tool catalogue	Count and rate cannot be mixed; missing is not zero; no causal, safety or desirability claims
Student 4	Professional-question objective plus two-tool catalogue	Questions for qualified review only; no legal, safety or compliance certification
Student 5	Case-summary objective plus four-tool catalogue	User text is untrusted; summary at most 120 words; three to five short actions; cite limitations

The feature objectives, IDs and tool catalogues are concrete prompt assets even where they are defined beside the route rather than in a separate template file. Their request limits range from three to four iterations, six to ten tool calls and a 120-second budget. Tool responses apply their own list and record bounds before any content enters the model context.

5.4 Agent workflow record

The strongest retained failure-recovery record is Feature 1 run `c7ca3ef0-c76f-47f5-bc42-acf03b20f8eb` from 22 August. It used prompt set `default.v4` under the Gemini development profile, performed three evidence calls over three iterations and joined the candidate release, failed ingestion run, quality results and accepted predecessor. It identified a zero-record draft and `stage_execution_failed`, preserved the 104-record accepted predecessor and recommended a bounded human-reviewed reprocess. It did not retry or publish.

Record field	Retained value
Objective	Diagnose a failed Feature 1 import using durable release and run evidence
Plan	Select release, run and predecessor inspection tools within the feature allowlist
Act	Execute three read-only evidence calls over HTTP
Observe	Candidate contained zero records, zero quality results and a failed stage; predecessor remained accepted with 104 records
Adapt	Explain the failure boundary and propose human-reviewed reprocessing without executing a write
Outcome	Succeeded after three iterations with evidence and safety note
Reproducible pointer	<code>student-1/AI_EVALUATION.md</code> and the durable run ID above

A later OpenAI run recorded in the report, `5a436835-e8e3-41f1-a00a-4c9a29b2454e`, used `default.v7`, Luna planning and Terra adaptation. The retained report account describes two Plan, Act, Observe and Adapt cycles using property search and property inspection. The final answer retained the verified address, accepted sale-history coverage, zero returned rows and a warning against inferring that the property never sold. This demonstrates the reported prompt refinement from operational recovery to typed provenance and ordered

multi-tool evidence. Its raw event export is not checked into the software baseline, so the run ID and narrative alone are not an independently replayable execution log. The earlier failed-import scenario has a linked [evaluation record](#).

5.5 AI-assisted software review record

The retained [26 July scaffold review](#) records an AI-assisted examination of the assignment, repository structure, database isolation, Compose, workflows and shared UI. It is a historical development artefact: its six-person roster, empty-Compose description and some proposed database patterns predate the approved five-feature architecture and must not be treated as current instructions.

Review area	Recorded finding	Current evidence to compare
Database design	Resolve store technology and ownership explicitly	Owned database APIs and migrations, with SQLite and PostgreSQL/PostGIS isolated within their feature boundaries
Implementation	Build shared styling and a unified home page	Shared HTMX entry point and the application screenshots in Section 8.4
Microservices	Populate the scaffold with container networking and service boundaries	Five individual diagrams, generated 21-service profile and architecture validator
DevOps	Align workflow names and configure branch triggers	The five student workflow files and retained successful runs in Section 7

This review provides retained development findings, while Section 5.4 provides application-loop execution evidence. The repository does not retain the complete original specialised review prompt and phase-by-phase terminal log for every student. The report therefore does not claim that the runtime diagnosis record alone proves every software-development review requirement.

6 Implementation summary

6.1 Shared platform

The shared platform supplies the HTMX entry point, common design tokens, feature registry, status and evidence views, reusable AI chat, mapping components, versioned contracts, consumer protocol, testkit, model registry, feature onboarding and generated deployment. The platform exposes only manifest-enabled routes, so placeholder folders do not appear as available features.

6.2 Feature capability and CRUD evidence

Feature	Visible CRUD	Principal public API	Seed or data evidence	AI path
Student 1	Source definitions and jobs; release review actions	/api/data-platform/v1	10 fixture properties; official runs include 5,190,134 G-NAF addresses and 7,335,504 PSI revisions	Failure diagnosis and property evidence
Student 2	Market-case create, list, detail, update and delete	/api/market-intelligence/v1	10 market cases and 26 sale observations	Explain a deterministic case summary
Student 3	Saved-comparison create, list, update and delete; browser bookmarks	/api/suburb-analytics/v1	10 suburbs, 10 overviews, 400 indicators, 38 amenities and 10 comparisons	Suburb snapshot and responsible comparison
Student 4	Site-review create, detail, edit, checklist and delete	/api/due-diligence/v1	10 reviews, 70 constraints and 50 building observations	Generate professional -verification questions
Student 5	Buyer cases, shortlist properties, notes and tasks	/api/buyer-workspaces/v1	10 cases, 12 properties, 11 notes and 12 tasks	Evidence-aware case summary and actions

6.3 Cross feature integration

Feature 2 validates property references and imports compatible sale releases through Feature 1 HTTP contracts. Feature 3 implements an idempotent consumer for accepted BOCSAR, schools and SEIFA products. Feature 4 searches Feature 1 for a verified property and uses its coordinates for a deterministic map. Feature 5 composes bounded evidence from Features 1, 2 and 4. It currently reports Feature 3 evidence as unavailable. None of these paths opens another feature's database.

7 DevOps and GitHub Actions

7.1 Pipeline architecture




Figure 15 GitHub Actions and container validation pipeline

Integration CI is the canonical source gate. It installs the locked workspace, runs formatting, linting, generation-drift checks, architecture and packaging validators, Mypy, Python tests, frontend tests, coverage and Compose configuration validation. Student workflows are path-filtered so each owner can build and test the assigned containers without repeating the entire gate.

7.2 Student workflow responsibilities

Workflow	Trigger and scope	Feature checks	Container evidence
student-1.yml	Pull requests and pushes affecting Student 1 or its shared boundary	Browser forms, source tests and official-source Compose plan	Builds Shared and Feature 1, starts bounded fixture stack, runs acquisition and boundary smoke
student-2.yml	Student 2 paths and manual dispatch	Ruff, architecture, generated deployment, Mypy, pytest and JavaScript	Builds three images, starts fixture stack, checks CRUD, seeds and frontend
student-3.yml	Student 3 paths and manual dispatch	Ruff, Mypy, branch-aware pytest and frontend behaviour	Builds three images, starts stack, checks frontend, API and seeded DB
student-4.yml	Student 4 paths and manual dispatch	Ruff, architecture, deployment, tool catalogue, Mypy, pytest and JavaScript	Builds four images, starts PostgreSQL slice, checks reads and seed counts
student-5.yml	Student 5 paths and manual dispatch	Ruff, Mypy, pytest, JavaScript and architecture validators	Builds three images, performs real CRUD, evidence and AI degradation smoke, restarts and proves persistence

Student 4's step name refers to CRUD but the shell body checks frontend GET, list GET and table counts; its create, update and delete paths are covered by API tests rather than that live container smoke.

7.3 Successful workflow evidence

Workflow	Evidence SHA	Successful run	Result
Integration CI	7d5350d	Run 33836061545	Canonical quality gate succeeded
Student 1 CI	f478ca6	Run 33785766339	Browser forms and integrated Feature 1 stack succeeded
Student 2 CI	5d650a0	Run 33631792022	Quality, images, Compose, CRUD, seeds and frontend succeeded
Student 3 CI	7d5350d	Run 33836061467	Quality and integrated-stack jobs succeeded
Student 4 CI	cfbbe60	Run 33776708519	Quality, four images, stack and seed smoke succeeded
Student 5 CI	cce16b2	Run 33810163715	Quality, CRUD, persistence and degradation smoke succeeded

The runs occur at different SHAs because path filters execute a student workflow only when its owned scope changes. The final baseline Integration and Student 3 runs succeeded after the last Feature 3 merge. The report does not imply that the five student jobs ran together at one SHA.

8 Testing and validation evidence

8.1 Local canonical quality gate

The command `uv run python scripts/check.py` was executed on Windows against baseline 7d5350d on 4 September 2026. Docker-dependent PostgreSQL component cases were skipped unless an opt-in disposable administrator URL was supplied; deterministic source tests did not require credentials or internet access.

Test group	Result	Coverage or scope
Format and Ruff	Passed	418 files formatted; lint clean
Contracts deployment architecture packaging models and tools	Passed	Five enabled feature catalogues and generated projections
Mypy	Passed	185 typed source files in canonical scope
Shared agent core and AI mode	602 passed 1 Windows symlink case skipped	90.36 percent branch-aware coverage
Student 1	587 passed 29 opt-in PostgreSQL cases skipped	72.60 percent coverage against 60 percent gate
Student 2	16 passed	77.10 percent coverage against 70 percent gate
Student 3	67 passed	90.25 percent branch-aware coverage against 80 percent gate
Student 4	60 passed	90.43 percent coverage against 85 percent gate
Student 5	136 passed	82.67 percent coverage against 80 percent gate
Frontend behaviour	Student 2 3 passed Student 3 12 Student 4 18 Student 5 21	Node tests plus shared and Feature 1 syntax and behaviour checks
Compose configuration	Passed	Production and development models with all five manifests

8.2 Endpoint and container evidence

Boundary	Verification
Public frontend to backend	Student workflow smokes request each feature frontend and public API health route
Backend to owned database API	Component tests use the real HTTP contract; container jobs check readiness and seeded reads
CRUD	Features 2 and 5 perform public container CRUD smokes; all five have API and persistence tests for create read update delete
Database ownership	Architecture validation checks credentials, volume mounts and forbidden imports
AI tools	Tool registry tests validate catalogue ownership, input and output schemas and safe dependency failure
Persistence	Migration tests prove idempotency; Student 5 restarts without deleting its volume and checks created and deleted state
Docker Compose	Every student workflow builds its images; the generated all-feature profile validates to 21 services

8.3 Non functional testing

Quality characteristic	Test or measure	Result and limit
Source scale	Complete PSI cached reprocess	7,079,728 staged and 6,667,588 accepted; 19 minute complete-source run evidence
Source integrity	Hash, byte and record verification	Content-addressed artefacts and exact manifest checks; corrupt or truncated input fails atomically
Cancellation	Runner and loader polling plus PostgreSQL statement cancellation	Active work observes durable cancellation and rolls back
Request size	Feature HTTP boundaries	Feature 4 256 KiB; Feature 5 1 MiB; import and tool limits are explicitly bounded
Concurrency	Mutable case and comparison resources	Features 2, 3 and 5 return 409 on stale versions
Accessibility	Static and behaviour tests	Labels, keyboard interaction, focus, live regions, table alternatives and 320 CSS-pixel floor
Degradation	Offline provider tests	CRUD and evidence remain available; AI returns an explicit unavailable state
Recovery	Import lease and receipt tests	Previous current release remains live after failure; replay is idempotent

The project does not claim a single end-user latency service-level objective for every feature. The source-scale benchmark measures real throughput and capacity; ordinary UI endpoints are bounded and smoke-tested, but a common percentile latency benchmark was not retained for Release 0.

8.4 Application screenshots




Figure 16 Shared entry point with five research areas




Figure 17 Feature 1 property discovery and attributed evidence




Figure 18 Feature 2 review of a deterministic market case




Figure 19 Feature 2 AI explanation of selected evidence

The screenshots are deterministic demonstration captures rather than proof of current official publisher facts.

9 Known issues and limitations

Limitation	Effect on Release 0 evidence
Feature 3 accepted official products were unavailable at the audited runtime	Its Release 0 views and AI remain clearly labelled deterministic fixtures
Feature 4 evidence and map layers are synthetic	The map illustrates evidence states around a property coordinate and must not be read as a parcel intersection or current planning certificate
Feature 5 queries only bounded candidate pages from Features 2 and 4	A matching record outside the first 25 can appear unavailable
Feature-specific live AI evidence is uneven	Feature 1 has durable run IDs and Feature 4 records a live evaluation; Features 2, 3 and 5 rely mainly on deterministic tool and degradation tests
Original software-review prompts and terminal phase logs are incomplete	Section 5.5 retains the historical AI-assisted review, but a per-student development-loop transcript is not present
No final local all-feature Docker execution log is retained	Per-feature CI stack runs and all-feature Compose validation are available; they do not prove one simultaneous local startup
No common percentile endpoint benchmark	Source-scale operations are measured, but ordinary feature endpoint latency is not reported as one cross-feature SLA

These limitations preserve the distinction between implemented behaviour, deterministic fixture evidence and official or production claims. They do not prevent the demonstrated local research workflows, but they constrain how results may be interpreted.

10 Contribution and commit record

10.1 Group contribution summary

Student	Principal contribution	Merged pull request evidence
Matthew	Shared platform, agent core, AI mode, contracts, deployment, Feature 1 data platform, HTMX migration and report foundation	71 merged PRs by MattShelton04 ; selected PRs 45, 51, 60 to 64, 66, 70, 73, 74 and 83
Burhan	Complete Feature 2 vertical slice, CI, walkthrough and friendly identifiers	PR 78 and PR 85
James	Feature 3 vertical slice, ingestion, comparisons, filters, accessibility and workflow fixes	PRs 81, 86, 90, 92 to 99
Michael	Feature 4 integration, CRUD, evidence, mapping, AI, CI and marking documentation	PRs 68, 69, 72, 75 to 80 and 87 to 89
Derek	Feature 5 store, cases, shortlist, notes, tasks, evidence, AI, runtime, CI and presentation refinement	PR 82 and PR 91

10.2 Selected immutable commits

Student	Date	Commit	Contribution
Matthew	29 August	2bf71fe	Migrated shared shell and Feature 1 Source CRUD to HTMX
Matthew	31 August	34f36f0	Added parallel read-only agent tool stages and v7 prompts
Matthew	2 September	0ac2c7f	Integrated ABS SEIFA 2021 into Feature 1
Burhan	2 September	5d650a0	Delivered Release 0 sales explorer and market cases
Burhan	3 September	bd6be66	Replaced internal identifiers with user-facing labels
James	3 September	a50dc70	Added the initial Suburb Analytics vertical slice
James	4 September	37e0100	Added comparison fixture data, units and loading state
James	4 September	7d5350d	Corrected filters without clearing the selected suburb
Michael	1 September	3b17c3a	Integrated the initial Due Diligence service slice
Michael	2 September	070163e	Added review edit and delete flows
Michael	4 September	36dbf71	Added bounded AI professional questions
Derek	3 September	7fc9c09	Delivered the Buyer Journey Release 0 feature
Derek	4 September	cce16b2	Refined evidence-aware AI presentation

The baseline shortlog attributes 61 commits to Matthew's two identities, 2 to Burhan, 11 to James, 11 to Michael and 2 squash commits to Derek. Derek's first pull request contains five logical commits before its squash merge. Pull request links provide the most useful review and integration record because they include workflow checks and the final merge state.

10.3 Attendance and demonstration

Checkpoint	Matthew	Burhan	James	Michael	Derek	Evidence
Feature implementation	Completed	Completed	Completed	Completed	Completed	Merged contribution record above
Recorded group demonstration	Completed	Completed	Completed	Completed	Completed	Published video linked on the cover
Week 6 class presentation	Attended	Attended	Attended	Attended	Attended	Group presentation and Q and A completed 4 September 2026

The published video is no longer than ten minutes and demonstrates the integrated product, each student feature's AI path, deployment and CI/CD. Agent-loop execution is documented in Section 5 in line with the later showcase instruction.

11 Rubric traceability

Criterion	Report evidence	Repository evidence	Readiness
1 Project Setup	Sections 1, 3 and 4	Root structure, five manifests, shared shell and Compose	Strong
2 Service Implementation	Sections 4 and 6	Five enabled service slices and health checks	Strong
3 AI Mode Integration	Section 5	Shared run API and five feature tool catalogues	Implemented with registered OpenAI selection
4 Agentic AI Workflow	Sections 5.2, 5.4 and 5.5	Agent core state machine, evaluation record and retained run IDs	Runtime implemented; development-review log gap disclosed
5 Prompt Engineering and Context	Section 5.3	Versioned shared prompts and feature objectives and schemas	Strong, feature assets are split between templates and routes
6 DevOps and GitHub Actions	Section 7	Five student workflows and Integration CI	Strong
7 Docker Compose Integration	Sections 4.3 and 8.2	One 21-service profile and successful feature stack jobs	Strong; no final local Docker rerun
8 Working Software	Section 6	CRUD tests, feature smokes, seeds and presentation	Strong
9 Technical Report	Sections 1 to 11	Five individual architectures and ERDs; tests, screenshots, logs and contributions	Covered; evidence gaps disclosed in Section 9
10 Project Demonstration	Section 10.3 and cover link	Published recording and completed Week 6 presentation	Complete when URL access is verified

Appendix A Evidence index

- [Root README](#)
- [Contributing guide](#)
- [Registered feature scope](#)
- [Shared platform design](#)
- [Integration contract](#)
- [Agent run state machine](#)
- [Feature 1 README](#)
- [Feature 1 marking evidence](#)
- [Feature 1 AI evaluation](#)
- [Feature 2 README](#)
- [Feature 3 README](#)
- [Feature 4 marking evidence](#)
- [Feature 5 implementation evidence](#)
- [Integration CI](#)
- [Student 1 CI](#)
- [Student 2 CI](#)
- [Student 3 CI](#)
- [Student 4 CI](#)
- [Student 5 CI](#)
- [Docker Compose model](#)

Appendix B Reproduction commands

```
uv sync --locked --all-packages --all-groups
uv run python scripts/check.py
uv run scripts/dev.py stack up
uv run scripts/dev.py stack status
uv run scripts/dev.py operator report
```

To regenerate this report from a clean checkout:

```
uv sync --locked --all-packages --all-groups
uv run python scripts/build_release0_report.py
```

When Mermaid sources change, install Node.js and run the builder with `--render-diagrams`. The pinned renderer refreshes PNGs and their content-hash manifest before generating the PDF. Normal builds use those checked-in assets offline and reject stale figures. Review the rendered pages after changing Markdown, figures or layout; a successful build alone is not a visual check.